

Certified Kubernetes Application Developer (CKAD)

Kubernetes 1.31 Edition




About the trainer

Brahim HAMDI

brahim.hamdi.consult@gmail.com

Consultant / Trainer
DevOps & Cloud Expert



Application Design and Build





Image and Container Management

Defining, building, and modifying container images

Container Terminology

The container and its runtime environment

Container

- Packages an application into a single unit of software including its runtime environment, application binary, dependencies, and configuration.

Container Runtime Engine

- A software component that can run containers on a host operating system.
- Example: [Docker Engine](#) or [containerd](#).

Container Orchestrator

- Uses a container runtime engine to instantiate a container.
- Automates and manages workload with features like scalability, networking.
- Example: [Kubernetes](#) or [Nomad](#)

Container Terminology

Building, running, and publishing a container image

Containerfile

- Spells out what needs to happen when the software is built in the form of a list of instructions.
- Synonymous with Dockerfile.

Container Image

- Packages an application into a single unit of software including its runtime environment and configuration.
- Built from the Containerfile.

Container Registry

- Shares container images in a central location and makes them available for consumption.
- Example: [Docker Hub](#) or [GCR](#)

Containerization Process

Building and publishing an image, running image in container




Image Building Instructions

A Dockerfile is a plain-text file with [instructions](#)

```
FROM      openjdk:11-jre-slim
WORKDIR   /app
COPY      target/java-hello-world-0.0.1.jar java-hello-world.jar
ENTRYPOINT ["java", "-jar", "/app/java-hello-world.jar"]
EXPOSE    8080
```


Building the Container Image

Provide name and tag of image + context directory

```
$ docker build -t java-hello-world:1.0.0 .
```

← Context directory

```
Sending build context to Docker daemon 8.32MB
```

```
Step 1/5 : FROM openjdk:11-jre-slim
```

```
---> 973c18dbf567
```

```
Step 2/5 : WORKDIR /app
```

```
---> Using cache
```

```
---> 31f9c5f2a019
```

```
...
```

```
Successfully built 3e9c22451a17
```

← Container Image ID

```
Successfully tagged java-hello-world:1.0.0
```

Listing Container Images

Render all local images including their name + hash

```
$ docker images
```

REPOSITORY	TAG	IMAGE ID	CREATED	SIZE
java-hello-world	1.0.0	3e9c22451a17	About a minute ago	213MB
openjdk	11-jre-slim	973c18dbf567	20 hours ago	204MB

Base image is downloaded automatically

Running the Container

You can interact with container once its running

```
$ docker run -d -p 8080:8080 java-hello-world:1.0.0  
b0ee04accf078ea7c...
```

```
$ docker logs b0ee04accf078ea7c...  
Hello World!
```

```
$ docker exec -it b0ee04accf078ea7c... bash  
root@b0ee04accf078ea7c:/# pwd  
/  
root@b0ee04accf078ea7c:/# exit
```

Tagging the Container Image

Name + tag need to conform to registry conventions

```
$ docker tag java-hello-world:1.0.0 bhamdi/java-hello-world:1.0.0
```

```
$ docker images
```

REPOSITORY	TAG	IMAGE ID	...
java-hello-world	1.0.0	3e9c22451a17	...
openjdk	11-jre-slim	973c18dbf567	...
bhamdi/java-hello-world	1.0.0	3e9c22451a17	...

↑
Tagged image

Publishing the Container Image

Authenticate against registry and push a tag

```
$ docker login --username=bhamdi
```

```
$ docker push bhamdi/java-hello-world:1.0.0
```

Saving a Container Image

Produces a backup TAR archive file with all parent layers

```
$ docker save -o java-hello-world.tar java-hello-world:1.0.0
```

```
$ ls
```

```
java-hello-world.tar
```



← openjdk:11-jre-slim

← java-hello-world:1.0.0

Loading a Container Image

Restores image from TAR archive file

```
$ docker load --input java-hello-world.tar
```

```
Loaded image: java-hello-world:1.0.0
```

```
$ docker images
```

REPOSITORY	TAG	IMAGE ID	...
java-hello-world	1.0.0	3e9c22451a17	...
openjdk	11-jre-slim	973c18dbf567	...

Lab 01

Defining, building, and
running a container image






Object Management

Kubernetes primitives, interaction with kubectl, imperative vs. declarative approach

Kubernetes Object Structure

General structure followed



General Sections

Each section has a specific function

- *API version*: Defines the structure and uses it to validate the correctness of the data.
- *Kind*: Defines the kind, e.g. a Pod or a Service.
- *Metadata*: Describes higher-level information about the object.
- *Spec*: Desired state of object.
- *Status*: Actual state of object. Empty if not created yet.

Using kubectl

Primary tool to interact with the Kubernetes clusters from the CLI

kubectl [command] [TYPE] [NAME] [flags]
 get pod app -o yaml

Imperative Object Management

Fast but requires detailed knowledge, no track record

```
$ kubectl create namespace ckad  
namespace/ckad created
```

Creates an object
e.g. a namespace

```
$ kubectl run nginx --image=nginx -n ckad  
pod/nginx created
```

Creates a Pod object
in a namespace

```
$ kubectl edit pod nginx -n ckad  
pod/nginx edited
```

Modifies the live
object e.g. a Pod

Declarative Object Management

Suitable for more elaborate changes, tracks changes

```
$ kubectl create -f nginx-deployment.yaml  
deployment.apps/nginx-deployment created
```

Creates object if it
doesn't exist yet

```
$ kubectl apply -f nginx-deployment.yaml  
deployment.apps/nginx-deployment configured
```

Creates object or update if
it already exists

```
$ kubectl delete -f nginx-deployment.yaml  
deployment.apps "nginx-deployment" deleted
```

Deletes object

Hybrid Approach

Generate YAML file with `kubectl` but make further edits

```
$ kubectl run nginx --image=nginx  
  --dry-run=client -o yaml > nginx-pod.yaml
```

```
$ vim nginx-pod.yaml
```

```
$ kubectl create -f nginx-pod.yaml  
pod/nginx created
```

← Capture YAML for
Pod object in file

← Create the object
from YAML file



Pods and Namespaces

Running and inspecting workload, organizing objects

What is a Pod?

Runs a process or application in container(s)

- The smallest, most basic deployable objects in Kubernetes.
- It contains one or more containers, executed by the configured container runtime.
- Kubernetes enhances the basic functionality of a container with features such as security, storage, and more.

Creating a Pod with Imperative Command

"Execute Hazelcast and configure it using parameters"

```
$ kubectl run hazelcast↵  
  --image=hazelcast/hazelcast↵  
  --port=5701↵  
  --env="DNS_DOMAIN=cluster"↵  
  --labels="app=hazelcast,env=prod"↵  
pod/hazelcast created
```

Pod YAML Manifest

"Execute Hazelcast and configure it using parameters"


```
apiVersion: v1
kind: Pod
metadata:
  name: hazelcast
spec:
  containers:
  - image: hazelcast/hazelcast
    name: hazelcast
    ports:
    - containerPort: 5701
```

← Container image

← Exposed container port(s)

Container Image Retrieval

Downloads container image from registry upon Pod creation



Listing Pods

Get high-level information on all Pods or a specific one

```
$ kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
hazelcast	1/1	Running	0	17s


```
$ kubectl get pods hazelcast
```

NAME	READY	STATUS	RESTARTS	AGE
hazelcast	1/1	Running	0	17s

← Specific Pod name

Pod Life Cycle Phases

Indicates operational status of Pod



Rendering Pod Details

Shows Pod configuration, events, and status

```
$ kubectl describe pods hazelcast
```

```
Name:                hazelcast
```

```
Namespace:           default
```

```
...
```

```
Status:              Running
```

```
IP:                  10.1.0.41
```

```
Containers:
```

```
...
```

```
Events:
```

```
...
```



Virtual IP address assigned to Pod to communicate with it

Accessing Container Logs

Renders logs produced by application or process running in a container

```
$ kubectl logs hazelcast
```

```
...
```

```
May 25, 2020 3:36:26 PM com.hazelcast.core.LifecycleService
```

```
INFO: [10.1.0.46]:5701 [dev] [4.0.1] [10.1.0.46]:5701 is STARTED
```

- You can stream the logs with the command line option `-f`. This option is helpful if you want to see logs in real time.
- You can still get back to the logs of the previous container by adding the `-p` command line option in case the container crashes or is restarted.

Executing a Command in a Container

Interactive exploration of a container, or seeing the direct effect of a command

```
$ kubectl exec -it hazelcast -- /bin/sh
# ...

$ kubectl exec hazelcast -- env
...
DNS_DOMAIN=cluster
```

- The command line option `-it` opens an interactive shell. This is useful if you want to inspect the configuration of your application or debug the current state of your application.
- For direct execution of a command in the container simply append to `--`.

Declaring Environment Variables

Define zero to many key-value pairs

```
apiVersion: v1
kind: Pod
metadata:
  name: spring-boot-app
spec:
  containers:
  - image: bhamdi/spring-boot-app:1.5.3
    name: spring-boot-app
    env:
    - name: SPRING_PROFILES_ACTIVE
      value: prod
    - name: VERSION
      value: '1.5.3'
```

Typical naming conventions
for keys are not enforced

Declaring Command and Args

Assign or overwrite container entrypoint and arguments

```
apiVersion: v1
kind: Pod
metadata:
  name: spring-boot-app
spec:
  containers:
  - image: bhamdi/spring-boot-app:1.5.3
    name: spring-boot-app
    command: ["/bin/sh"]
    args: ["-c", "while true; do date; sleep 10; done"]
```

Deleting a Pod

Graceful deletion of workload can take up to 30 seconds

```
$ kubectl delete pod hazelcast  
pod "hazelcast" deleted
```

← Imperative approach

```
$ kubectl delete -f pod.yaml  
pod "hazelcast" deleted
```

← Declarative approach Force-kills

```
$ kubectl delete -f pod.yaml --now  
pod "hazelcast" deleted
```

← Pod without
graceful deletion of workload

What is a Namespace?

Organizes and isolates Kubernetes object that belong together

- An API construct to avoid naming collisions and represent a scope for object names. A good use case for namespaces is to isolate the objects by team or responsibility.
- The `default` namespace hosts object that haven't been assigned to an explicit namespace.
- Namespaces starting with the prefix `kube-` are not considered end user-namespaces. They have been created by the Kubernetes system. You will not have to interact with them as an application developer.
- Reference a namespace with the command line option `--namespace` or `-n`.

Creating a Namespace with Imperative Command

"Organize objects in namespace called `code-red`"

```
$ kubectl create namespace code-red
```

```
namespace/code-red created
```

```
$ kubectl run nginx --image=nginx --restart=Never -n code-red
```

```
pod/nginx created
```

```
$ kubectl get pods -n code-red
```

NAME	READY	STATUS	RESTARTS	AGE
nginx	1/1	Running	0	13s

Listing Namespaces

Includes `default`, `Kubernetes-internal`, and custom namespaces

```
$ kubectl get namespaces
```

NAME	STATUS	AGE
default	Active	157d
kube-node-lease	Active	157d
kube-public	Active	157d
kube-system	Active	157d

Namespace YAML Manifest

"Organize objects in namespace called `code-red`"

```
apiVersion: v1
kind: Namespace
metadata:
  name: code-red
```


Deleting a Namespace

Cascade-deletes containing objects

```
$ kubectl delete namespace code-red  
namespace "code-red" deleted
```

```
$ kubectl get pods -n code-red  
No resources found in code-red namespace.
```

Lab 02

Creating, inspecting, and
interacting with a Pod in a
Namespace





Jobs and CronJobs

One-time operations and scheduled operations


What is a Job?

Kubernetes primitive for executing a one-time operation

- It runs functionality until a specified number of completions has been reached.
- The actual work managed by a Job is still running inside of a Pod. Therefore, you can think of a Job as a higher-level coordination instance for Pods executing the workload.
- Upon completion of a Job and its Pods, Kubernetes does not automatically delete the objects - they will stay until they're explicitly deleted. Keeping those objects helps with debugging the command run inside of the Pod and gives you a chance to inspect the logs.

Job Object Hierarchy

A Job manages Pods to run the workload



Creating a Job with Imperative Command

"Increment a counter and render its value on the terminal"

```
$ kubectl create job counter
--image=nginx:1.24.0
-- /bin/sh -c 'counter=0; while [ $counter -lt 3 ]; do
    counter=$((counter+1)); echo "$counter"; sleep 3; done;'
job.batch/counter created
```

Job YAML Manifest

"Increment a counter and render its value on the terminal"

```
apiVersion: batch/v1
kind: Job
metadata:
  name: counter
spec:
  completions: 1
  parallelism: 1
  backoffLimit: 6
  template:
    spec:
      restartPolicy: OnFailure
    containers:
    - args:
      - /bin/sh
      - -c
      - ...
      image: nginx:1.24.0
      name: counter
```

Define # of successful completions and whether task should be run in parallel

How many times do we try before Job is marked failed?

Restart Pod upon failure or start a new Pod

Inspecting a Job

Job is completed when configured # of executions has been reached

```
$ kubectl get jobs
```

NAME	COMPLETIONS	DURATION	AGE
counter	0/1	13s	13s

```
$ kubectl get jobs
```

NAME	COMPLETIONS	DURATION	AGE
counter	1/1	15s	19s

```
$ kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
counter-z6kdj	0/1	Completed	0	51s

Inspecting Pod Logs

Workload can be inspected during and after execution

```
$ kubectl logs counter-z6kdj
```

```
1
```

```
2
```

```
3
```

Job Operation Types

Execution runtime behavior of workload can be fine-tuned

- The default behavior of a Job is to run the workload in a single Pod and expect one successful completion (non-parallel Job).
- The attribute `spec.completions` controls the number of required successful completion.
- The attribute `spec.parallelism` allows for executing the workload by multiple pods in parallel.

Job Operation Types

Pick combination of completions and parallelism based on workload use case

Type	Completions	Parallelism	Description
Non-parallel with one completion count	1	1	Completes as soon as its Pod terminates successfully.
Parallel with a fixed completion count	≥ 1	≥ 1	Completes when specified number of tasks finish successfully.
Parallel with worker queue	unset	≥ 1	Completes when at least one Pod has terminated successfully and all Pods are terminated.

Restart Behavior


Re-running workload if execution is unsuccessful

- The attribute `spec.backoffLimit` determines the number of retries a Job attempts to successfully complete the workload until the executed command finishes with an exit code 0. The default is 6, which means it will execute the workload 6 times before the Job is considered unsuccessful.
- The attribute `spec.template.spec.restartPolicy` needs to be declared explicitly. The restart policy of a Job can only be `OnFailure` or `Never`.

Restarting a Container on Failure

Restart container if it exits with a non-zero exit code


```
spec.template.spec.restartPolicy: OnFailure
```



Starting a New Pod on Failure

The policy does not restart the container upon a failure, it starts a new Pod.


```
spec.template.spec.restartPolicy: Never
```



What is a CronJob?


Kubernetes primitive for executing workload periodically based a schedule

- A CronJob is essentially a Job with similar characteristics.
- The schedule can be defined with a cron-expression you may already know from Unix cron jobs.



CronJob Object Hierarchy

A CronJob manages Jobs which delegate the workload to Pods



Creating a CronJob with Imperative Command

"Render the current date on standard output every hour"

```
$ kubectl create cronjob↵  
current-date↵  
--schedule="* * * * * "↵  
--image=nginx:1.24.0  
-- /bin/sh -c 'echo "Current date: $(date) "'  
cronjob.batch/current-date created
```

CronJob YAML Manifest

"Increment a counter and render its value on the terminal"

```
apiVersion: batch/v1
kind: CronJob
metadata:
  name: counter
spec:
  schedule: "*/* * * * *"
  jobTemplate:
    spec:
      template:
        spec:
          restartPolicy: Never
          containers:
            - args:
              - /bin/sh
              - -c
              - ...
              image: nginx:1.24.0
              name: counter
```

← The crontab expression used to run CronJob periodically

← Run in a new Pod

Inspecting a CronJob

Keeps a history of failed and successful Pods

```
$ kubectl get cronjobs
```

NAME	SCHEDULE	SUSPEND	ACTIVE	LAST SCHEDULE	AGE
counter	*/1 * * * *	False	0	26s	1h

```
$ kubectl get jobs --watch
```

NAME	COMPLETIONS	DURATION	AGE
counter-1557334380	1/1	3s	2m24s
counter-1557334440	1/1	3s	84s
counter-1557334500	1/1	3s	24s

Configuring Retained History

Being able to know what happened during workload execution retroactively

- Even after completing a task in a Pod controlled by a CronJob, it will not be deleted automatically. Keeping a historical record of Pods can be tremendously helpful for troubleshooting failed workloads or inspecting the logs.
- The attribute `spec.successfulJobsHistoryLimit` configures the number of successful executions. Defaults to 3 successful Jobs.
- The attribute `spec.failedJobsHistoryLimit` configures the number of failed executions. Defaults to 1 failed Job.

Lab 03

Creating and inspecting
a periodic operation
using a CronJob





Volumes

Ephemeral and persistent storage for Pods

What is a Volume?

A directory, possibly with some data in it accessible to the containers in a Pod

- Each container manages files in a temporary file system. Data written to the file system is lost when the container is restarted.
- *Ephemeral Volumes* exist for the lifespan of a Pod. They are useful if you want to share data between multiple containers running in the Pod or caching of data.
- *Persistent Volumes* preserve data beyond the lifespan of a Pod. They are a good option for applications that require data to exist longer, e.g. in the form of storage for a database-driven application.
- Define the Volume type with `spec.volumes[]` and then reference it in a container with `spec.containers[].volume.volumeMounts`.


Volume Types

Determines the medium that backs the Volume and its runtime behavior

Type	Description
<code>emptyDir</code>	Empty directory in Pod with read/write access. Only persisted for the lifespan of a Pod.
<code>hostPath</code>	File or directory from the host node's filesystem.
<code>configMap</code> , <code>secret</code>	Provides a way to inject configuration data.
<code>nfs</code>	An existing NFS (Network File System) share. Preserves data after Pod restart.
<code>persistentVolumeClaim</code>	Claims a Persistent Volume.

Ephemeral Volume

Useful for sharing data between containers, or for caching data



Defining an Ephemeral Volume

Provide type and assign mount path per container

```
apiVersion: v1
kind: Pod
metadata:
  name: my-container
spec:
  volumes:
  - name: logs-volume
    emptyDir: {}
  containers:
  - image: nginx
    name: my-container
    volumeMounts:
    - mountPath: /var/logs
      name: logs-volume
```

← Define Volume name and type

← Specify mount path in container

Using a Volume


You can shell into container and navigate to mount path

```
$ kubectl create -f pod-with-vol.yaml
pod/my-container created

$ kubectl exec -it my-container -- /bin/sh
# cd /var/logs
# pwd
/var/logs
# touch app-logs.txt
# ls
app-logs.txt
```

Persistent Volume

Persist data that outlives a Pod, node, or cluster restart



Static vs. Dynamic Provisioning

PersistentVolume object is created manually or automatically

- *Static Provisioning:* Storage device needs to be created first. The PersistentVolume object references the storage device and needs to be created manually.
- *Dynamic Provisioning:* The PersistentVolumeClaim object references a storage class. The PersistentVolume object is created automatically.
- *Storage Class:* Object that knows how to provision a storage device with specific performance requirements.

Defining a PersistentVolume

Define storage capacity, access mode, and host path

```
apiVersion: v1
kind: PersistentVolume
metadata:
  name: db-pv
spec:
  capacity:
    storage: 1Gi
  accessModes:
    - ReadWriteOnce
  hostPath:
    path: /data/db
```

Access Mode

Defines read/write capabilities of Volume

Type	Description
<code>ReadWriteOnce</code>	Read-write access by a single node.
<code>ReadOnlyMany</code>	Read-only access by many nodes.
<code>ReadWriteMany</code>	Read-write access by many nodes.
<code>ReadWriteOncePod</code>	Read/write access mounted by a single Pod.

Reclaim Policy

What should happen to PersistentVolume when Claim is deleted?

Type	Description
Retain	Default. When PVC is deleted, PV is “released” and can be reclaimed.
Delete	Deletion removes PV and associated storage.
Recycle	<i>Deprecated.</i> Use dynamic binding instead.

Creating a PersistentVolume

Summarized most important configuration and status

```
$ kubectl create -f db-pv.yaml
persistentvolume/db-pv created
```

```
$ kubectl get pv db-pv
```

NAME	CAPACITY	ACCESS	MODES	RECLAIM	POLICY	STATUS	...
db-pv	1Gi	RWO		Retain		Available	...



The object hasn't been consumed by a claim yet

Defining a PersistentVolumeClaim

Will bind to a PersistentVolume based on resource request

```
apiVersion: v1
kind: PersistentVolumeClaim
metadata:
  name: db-pvc
spec:
  accessModes:
    - ReadWriteOnce
  resources:
    requests:
      storage: 256Mi
      storageClassName: ""
```

Creating a PersistentVolumeClaim

Binds to PersistentVolume if requirements can be fulfilled

```
$ kubectl create -f db-pvc.yaml
persistentvolumeclaim/db-pvc created
```

```
$ kubectl get pvc db-pvc
```

NAME	STATUS	VOLUME	CAPACITY	ACCESS MODES	STORAGECLASS	AGE
db-pvc	Bound	pvc	512m	RWO		7s



Binding to the PersistentVolume
object was successful

Mounting a PersistentVolumeClaim in a Pod

Use Volume type `persistentVolumeClaim`

```
apiVersion: v1
kind: Pod
metadata:
  name: app-consuming-pvc
spec:
  volumes:
    - name: app-storage
      persistentVolumeClaim:
        claimName: db-pvc
  containers:
    - image: alpine
      ...
    volumeMounts:
      - mountPath: "/mnt/data"
        name: app-storage
```

← Reference the Volume by claim name

Assigning a StorageClass for Dynamic Provisioning

Creates PersistentVolume object automatically via storage class

```
apiVersion: storage.k8s.io/v1
kind: StorageClass
metadata:
  name: aws-ebs
provisioner: kubernetes.io/aws-ebs
```

```
apiVersion: v1
kind: PersistentVolumeClaim
metadata:
  name: pvc
spec:
  accessModes:
    - ReadWriteMany
  resources:
    requests:
      storage: 256Mi
  storageClassName: aws-ebs
```

Dynamic Creation of PersistentVolume Object

Adds hash to the name of the object

```
$ kubectl get pv,pvc
```

NAME		STORAGECLASS
persistentvolume/	db-pvc-b820b919-f7f7-4c74-9212-ef259d421734	aws-ebs

NAME	...	STORAGECLASS
persistentvolumeclaim/db-pvc		aws-ebs

Created by storage
class provisioner

Lab 04

Creating and using a PersistentVolume using static provisioning





Multi-Container Pods

Defining and interacting with multiple containers, design patterns

Defining Multiple Containers in a Pod

There are warranted use cases for wanting to run more than a single container

- The general guideline is to operate a microservice in a single container per Pod which promotes a decentralized, decoupled, and distributed architecture.
- Specific use cases call for running multiple containers per Pod. For example, you may want to provide helper functionality that runs alongside the application. In other cases, you may want to initialize your Pod by executing setup scripts, commands, or any other kind of preconfiguration procedure before the application container should start.

Design Patterns

Emerged from Pod requirements

Init Container

- Provide initialization logic concerns to be run before the main application even starts.
- Example: Downloading a configuration files required by application.

Adapter

- Transforms the output produced by the application to make it consumable in the format needed by another part of the system.
- Example: Massaging log data.

Sidecar


- The sidecars are not part of the main traffic or API of the primary application and operate asynchronously.
- Example: Watcher capabilities.

Ambassador

- Provides a proxy for communicating with external services to hide and/or abstract the complexity.
- Example: Rate-limiting functionality for HTTP(S) calls to an external service.

Init Container Life Cycle

Initialization logic before main application containers



Defining an Init Container

Declared adjacent to main application containers with `spec.initContainers`

```
apiVersion: v1
kind: Pod
metadata:
  name: multi-container
spec:
  initContainers:
  - image: init:3.2.1
    name: app-initializer
  containers:
  - image: nginx
    name: web-server
```

← Uses the same attributes as
main application containers

Creating a Pod with an Init Container

Init container execution renders in status column

```
$ kubectl create -f init.yaml
```

```
pod/business-app created
```

```
$ kubectl get pod business-app
```


NAME	READY	STATUS	RESTARTS	AGE
business-app	0/1	Init:0/1	0	2s

```
$ kubectl get pod business-app
```

NAME	READY	STATUS	RESTARTS	AGE
business-app	1/1	Running	0	8s

Interacting with a Container

Run `kubectl` with `--container` or `-c` CLI option



Shelling into a Container

Defaults to main application container if not explicit

```
$ kubectl exec -it -c web-server business-app -- /bin/sh
# exit
```

```
$ kubectl exec -it business-app -- /bin/sh
Defaulted container "web-server" out of: web-server, ←
app-initializer (init)
# exit
```

Defining More Than One Main Application Container

Add another container definition under `spec.containers`

```
apiVersion: v1
kind: Pod
metadata:
  name: multi-container
spec:
  containers:
    - image: nginx:1.13.0
      name: web-server
    - image: transformer:1.0.0
      name: helper
```

Defines a second container
running alongside the
application container

Creating a Pod with Multiple Application Containers

The ready column indicates the number of started containers

```
$ kubectl create -f multi-container.yaml
```


```
pod/multi-container created
```

```
$ kubectl get pod business-app
```

NAME	READY	STATUS	RESTARTS	AGE
multi-container	2/2	Running	0	2s


Sidecar Pattern

"Send a notification of error log contains entries"




Adapter Pattern

"Transform log data to JSON to make it consumable by external system"



Ambassador Pattern

"OAuth authentication handled by network communication proxy"



Lab 05

Implementing a mutli-container pod



Labels and Annotations

Label assignment and selection, providing meta information to objects


What is a Label?

Essential to querying, filtering and sorting Kubernetes objects

- Key-value pairs assigned to Kubernetes objects with the imperative `label` command or using the `metadata.labels[]` attribute.
- Not meant for elaborate, multi-word descriptions of its functionality.
- Kubernetes limits the length of a label to a maximum of 63 characters and a range of allowed alphanumeric and separator characters.

Example Labels

Three different Pods with each three key-value pairs



Assigning Labels with Imperative Approach

The `label` command lets you add or remove one or many labels

```
$ kubectl label pod nginx tier=backend env=prod app=miracle
pod/nginx labeled
```

← Declares three label key-value pairs

```
$ kubectl get pods --show-labels
NAME    ... LABELS
pod1    ... tier=backend,env=prod,app=miracle
```

```
$ kubectl label pod nginx tier-
pod/nginx labeled
```

← Removes the label with key `tier`

```
$ kubectl get pods --show-labels
NAME    ... LABELS
pod1    ... env=prod,app=miracle
```

Assigning Labels in a YAML Manifest

Can be provided for a new object or modified for a live object

```
apiVersion: v1
kind: Pod
metadata:
  name: pod1
  labels:
    tier: backend
    env: prod
    app: miracle
spec:
  ...
```

← Declares three label
key-value pairs


Recommended Labels

Naming convention starts with `app.kubernetes.io`

```
apiVersion: apps/v1
kind: Deployment
metadata:
  labels:
    app.kubernetes.io/name: wordpress
    app.kubernetes.io/instance: wordpress-abcxyz
    app.kubernetes.io/version: "4.9.4"
    app.kubernetes.io/managed-by: helm
    app.kubernetes.io/component: server
    app.kubernetes.io/part-of: wordpress
  ...
```

Selecting Labels

Label selector uses a set of criteria to query for objects



Label Selection from the CLI

You can specify equality-based and set-based requirements

```
$ kubectl get pods -l tier=frontend,env=dev --show-labels
```

```
NAME    ... LABELS
```

```
pod2    ... app=crawler,env=dev,tier=frontend
```

Boolean AND
expression

```
$ kubectl get pods -l version --show-labels
```

```
NAME    ... LABELS
```

```
pod3    ... app=crawler,env=staging,version=v2.1
```

Key-only
expression

```
$ kubectl get pods -l 'tier in (frontend,backend),env=dev' --show-labels
```

```
NAME    ... LABELS
```

```
pod2    ... app=crawler,env=dev,tier=frontend
```

Set-based expression
follows Boolean OR

Label Selection in YAML

Certain objects provide label selection via their API

```
apiVersion: networking.k8s.io/v1
kind: NetworkPolicy
metadata:
  name: frontend-network-policy
spec:
  podSelector:
    matchLabels:
      tier: frontend
  ...
```


What is an Annotation?

Descriptive metadata without the ability to be queryable

- Key-value pairs for providing descriptive, human-readable metadata assigned to Kubernetes objects with the imperative `annotation` command or using the `metadata.annotations[]` attribute.
- Cannot be used for querying or selecting objects.
- Make sure to put the value of an annotation into single- or double-quotes if it contains special characters or spaces.

Example Annotations

A Pod with three key-value pairs



Assigning Annotations with Imperative Approach

The `annotation` command lets you add or remove one or many annotations

```
$ kubectl annotate pod nginx commit='866a8dc' branch='bm/bugfix'  
pod/nginx annotated
```

← Declares three
annotation
key-value pairs

```
$ kubectl describe pods my-pod  
Annotations:  branch: bm/bugfix  
              commit: 866a8dc
```

```
$ kubectl annotate pod nginx commit-  
pod/nginx annotated
```

← Removes the
annotation with key
commit

```
$ kubectl describe pods my-pod  
Annotations:  branch: bm/bugfix
```

Assigning Annotations in a YAML Manifest

Can be provided for a new object or modified for a live object

```
apiVersion: v1
kind: Pod
metadata:
  name: pod1
  annotations:
    commit: 866a8dc
    author: 'Benjamin Muschko'
    branch: 'bm/bugfix'
spec:
  ...
```

← Declares three annotation
key-value pairs

Well-Known Annotations

Reserved annotations that may influence runtime treatment

```
apiVersion: v1
kind: Namespace
metadata:
  name: secured
  annotations:
    kubernetes.io/description: "Manages objects for business app."
    pod-security.kubernetes.io/warn: "baseline"
```

Enforces Pod
Security Standards
at the namespace
level

Lab 06

Using Labels and Annotations



Application Deployment




What is a Deployment?

Scaling and replication features for a set of Pods

- Controls a predefined number of Pods with the same configuration, so-called *replicas*.
- The number of replicas can be scaled up or down to fulfill load requirements.
- Updates to the replica configuration can be updated easily and is rolled out automatically.

Example Deployment

A Deployment that controls three replicas



Creating a Deployment with Imperative Approach

Creates objects for the Deployment, ReplicaSet and Pod(s)

```
$ kubectl create deployment↵  
my-deploy↵  
--image=nginx:1.14.2↵  
--replicas=3 ↵  
deployment.apps/my-deploy created
```

The default number of replicas is 1 if the parameter wasn't provided

Deployment Template Attributes

Replica configuration is defined under `spec.template`

```
apiVersion: apps/v1
kind: Deployment
metadata:
  labels:
    app: my-deploy
  name: my-deploy
spec:
  replicas: 3
  selector:
    matchLabels:
      tier: backend
  template:
    metadata:
      labels:
        tier: backend
    spec:
      containers:
        - image: nginx:1.14.2
          name: nginx
```

Deployment

Template spec attributes
are the same as for a
Pod spec



Pod


```
apiVersion: v1
kind: Pod
metadata:
  name: nginx
spec:
  containers:
    - image: nginx:1.14.2
      name: nginx
```

Deployment YAML Manifest

Creates an error if label selector and template labels do not match

```
apiVersion: apps/v1
kind: Deployment
metadata:
  labels:
    app: my-deploy
    name: my-deploy
spec:
  replicas: 3
  selector:
    matchLabels:
      tier: backend
  template:
    metadata:
      labels:
        tier: backend
    spec:
      containers:
        - image: nginx:1.14.2
          name: nginx
```

Label selector and
template label assignment
have to match



Listing Deployments

ReplicaSet and Pods can be identified by name prefix

```
$ kubectl get deployments,pods,replicasets
```

NAME	READY	UP-TO-DATE	AVAILABLE	AGE
deployment.apps/my-deploy	1/1	1	1	7m56s

NAME	READY	STATUS	RESTARTS	AGE
pod/my-deploy-8448c488b5-mzx5g	1/1	Running	0	7m56s

NAME	DESIRED	CURRENT	READY	AGE
replicaset.apps/my-deploy-8448c488b5	1	1	1	7m56s

Rendering Deployment Details

Object details show relationship between parent and child

```
$ kubectl describe deployment.apps/my-deploy
```

```
Replicas:          1 desired | 1 updated | 1 total | 1 available |  
                  0 unavailable
```

```
NewReplicaSet:    my-deploy-8448c488b5 (1/1 replicas created)
```

```
$ kubectl describe replicaset.apps/my-deploy-8448c488b5
```


```
Controlled By:    Deployment/my-deploy
```

```
$ kubectl describe pod/my-deploy-8448c488b5-mzx5g
```

```
Controlled By:    ReplicaSet/my-deploy-8448c488b5
```

Roll Out and Roll Back Feature

"Look ma, shiny new features. Let's deploy them to production!"



Rolling Out a New Revision

The rollout history keeps track of changes

```
$ kubectl rollout history deployment my-deploy
```

```
deployment.apps/my-deploy  REVISION  CHANGE-  
CAUSE
```

```
1          <none>
```

```
$ kubectl set image deployment my-deploy nginx=nginx:1.19.2
```

```
deployment.apps/my-deploy image updated
```

```
$ kubectl rollout history deployment my-deploy
```

```
deployment.apps/my-deploy  REVISION  CHANGE-  
CAUSE
```

```
1          <none>
```

```
2          <none>
```

← Changes the
assigned image
for Pod template

Rolling Back to Previous Revision

You can revert to any revision available in the history

```
$ kubectl rollout undo deployment my-deploy --to-revision=1
deployment.apps/my-deploy rolled back
```

```
$ kubectl rollout history deployment my-deploy
deployment.apps/my-deploy  REVISION  CHANGE-CAUSE
2                          <none>
3                          <none>
```

↑
Kubernetes deduplicates a revision
if it points to the same changes

Setting a Change Cause for a Revision

Tracked by annotation with key `kubernetes.io/change-cause`

```
$ kubectl rollout history deployment my-deploy
deployment.apps/my-deploy  REVISION  CHANGE-
CAUSE
2                          <none>
3                          <none>
```


```
$ kubectl annotate deployment my-deploy kubernetes.io/change-cause="image updated to 1.16.1"
deployment.apps/my-deploy annotated
```

```
$ kubectl rollout history deployment my-deploy
deployment.apps/my-deploy  REVISION  CHANGE-
CAUSE
2                          <none>
3                          image updated to 1.16.1
```

← What changed
with this revision?

Manually Scaling a Deployment

"We measured performance under load. Scale up to exactly x number of Pods."



Changing the Number of Replicas

Use scale command or change `spec.replicas` attribute in live object

```
$ kubectl scale deployment my-deploy --replicas=5  
deployment.apps/my-deploy scaled
```

```
$ kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
my-deploy-8448c488b5-5f5tg	1/1	Running	0	44s
my-deploy-8448c488b5-9xplx	1/1	Running	0	44s
my-deploy-8448c488b5-d8q4t	1/1	Running	0	44s
my-deploy-8448c488b5-f5kkm	1/1	Running	0	44s
my-deploy-8448c488b5-mzx5g	1/1	Running	0	3d19h

Lab 07

Performing rolling updates and manually scaling a Deployment




Types of Autoscalers

Scaling decisions are made based on metrics

- *Horizontal Pod Autoscaler (HPA)*: A standard feature of Kubernetes that scales the number of Pod replicas based on CPU and memory thresholds.
- *Vertical Pod Autoscaler (VPA)*: Scales the CPU and memory allocation for existing Pods based on historic metrics. Supported by a cloud provider as an add-on or needs to be installed manually.
- Both autoscalers use the [metrics server](#). You need to install the component. A Pod should also set the resource requests and limits.

Horizontal Pod Autoscaler (HPA)


Kubernetes primitive that reaches into metrics server



Autoscaling a Deployment by CPU

"Auto-scale based on average CPU utilization across all replicas."

Maximum, average CPU utilization: 80%



Horizontal Pod Autoscaler YAML Manifest

Creates an error if label selector and template labels do not match

```
apiVersion: autoscaling/v2
kind: HorizontalPodAutoscaler
metadata:
  name: app-cache
spec:
```

```
  scaleTargetRef:
    apiVersion: apps/v1
    kind: Deployment
    name: app-cache
```

```
  minReplicas: 3
```

```
  maxReplicas: 5
```

```
  metrics:
```

```
- type: Resource
  resource:
    name: cpu
    target:
      type: Utilization
      averageUtilization: 80
```

← The scaling target,
a Deployment

← One or many resource types that
the thresholds should apply to

Inspecting the Horizontal Pod Autoscaler

Use `scale` command or change `spec.replicas` attribute in live object

```
$ kubectl get hpa
```

NAME	REFERENCE	TARGETS	MINPODS	MAXPODS	REPLICAS
app-cache	Deployment/app-cache	15%/80%	3	5	AGE 3
2m14s					

↑
Renders current resource utilization
and compares them with thresholds

Lab 08

Creating a Horizontal Pod Autoscaler for a Deployment




Common Deployment Strategies

Choosing the right deployment procedure depends on the needs

- *Rolling Update*: Release a new version on a rolling update fashion, one after the other.
- *Recreate*: Terminate the old version and release the new one.
- *Blue/green*: Release a new version alongside the old version then switch traffic.
- *Canary*: Release a new version to a subset of users, then proceed to a full rollout.

Rolling Update Deployment Strategy

Secondary ReplicaSet is created with new version of application



Rolling Update Deployment Strategy

Runtime behavior can be fine-tuned


```
spec:  
  replicas: 3  
  strategy:  
    type: RollingUpdate  
    rollingUpdate:  
      maxSurge: 2  
      maxUnavailable: 0
```

← Determines how update
should be performed

- *Pros:* New version is slowly distributed over time, no downtime
- *Cons:* Breaking API changes can make consumers incompatible

Recreate Deployment Strategy

Terminates all the running instances then recreate them



Recreate Deployment Strategy

No configuration options available


```
spec:  
  replicas: 3  
  strategy:  
    type: Recreate
```

← Set the strategy explicitly

- *Pros:* Good option for development environments as everything is renewed at once
- *Cons:* Can cause downtime while applications are starting up

Blue/Green Deployment Strategy

Deployment object with new version is created alongside



Blue/Green Deployment Strategy

Requires a duplicate set of objects


```
apiVersion: apps/v1
kind: Deployment
metadata:
  labels:
    app: my-deploy
    version: 1.0.0
name: blue
```

```
apiVersion: apps/v1
kind: Deployment
metadata:
  labels:
    app: my-deploy
    version: 1.1.0
name: green
```

- *Pros:* No downtime, traffic can be routed when ready
- *Cons:* Resource duplication, configuration of network routing

Canary Deployment Strategy

Two Deployment objects with different traffic distribution



Canary Deployment Strategy

The new Deployment initially starts with a smaller number of replicas

```
spec:  
  replicas: 3
```

```
spec:  
  replicas: 1
```

- *Pros:* New version released to subset of users, A/B testing of features and performance
- *Cons:* May require a load balancer for fine-grained distribution

Lab 09

Implementing the blue-green deployment strategy





Helm

Deploying an application defined by a set of YAML manifests

What is Helm?

Package manager and templating engine for a set of manifests

- The artifact produced by the Helm executable is a so-called *chart file* bundling the manifests that comprise the API resources of an application.
- Chart files can be distributed to a chart repository e.g. [central chart repository](#) and consumed from there (e.g., for running Grafana or PostgreSQL).
- At runtime, it replaces placeholders in YAML template files with actual, end-user defined values.




Finding a Chart

Searches the Artifact Hub for a chart with a matching name

```
$ helm search hub jenkins
```

URL	CHART	VERSION	APP
https://artifacthub.io/packages/helm/bitnami/jenkins	11.0.6	2.361.2	
...			



Installing a Chart

You may have to add an external repository that hosts chart file

```
$ helm repo add bitnami https://charts.bitnami.com/bitnami
"bitnami" has been added to your repositories
```

```
$ helm install jenkins bitnami/jenkins
```

```
...
```

```
CHART NAME: jenkins
```

```
CHART VERSION: 11.0.6
```

```
APP VERSION: 2.361.2
```

```
** Please be patient while the chart is being deployed **
```

Listing Installed Charts

Information about chart details

```
$ helm list
```

NAME	NAMESPACE	REVISION	...	STATUS	CHART	APP VERSION
jenkins	default	1		deployed	jenkins-11.0.6	2.361.2

```
$ kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
pod/bitnami-jenkins-764b44cc99-tw6f6	1/1	Running	0	7m34s
...				

Uninstalling a Chart

Deletes Kubernetes objects controlled by chart

```
$ helm uninstall jenkins
```

```
release "jenkins" uninstalled
```

```
$ kubectl get pods
```

```
No resources found in default namespace.
```

Standard Chart Structure

Predefined directory structure with mandatory `Chart.yaml` and `values.yaml`

```
$ tree
```

```
.
├── Chart.yaml
├── templates
│   ├── web-app-pod-template.yaml
│   └── web-app-service-template.yaml
└── values.yaml
```

Chart File

Describes the meta information of the chart (e.g., name and version)

Chart.yaml

```
apiVersion: 1.0.0  
name: web-app  
version: 2.5.4
```

Values File

Key-value pairs used at runtime to replace placeholders in the YAML manifests

values.yaml

```
db_host: mysql-service
db_user: root
db_password: password
service_port: 3000
```

Additional [built-in objects](#) are available for use.

Template File

YAML manifest that can define placeholders using double curly braces

web-app-pod-template.yaml

```
apiVersion: v1
kind: Pod
metadata:
  name: web-app
spec:
  containers:
  - image: bmuschko/web-app:1.0.1
    name: web-app
    env:
    - name: DB_HOST
      value: {{ .Values.db_host }}
    - name: DB_USER
      value: {{ .Values.db_user }}
    - name: DB_PASSWORD
      value: {{ .Values.db_password }}
    ...
```

Previewing Final Templates

Replaces placeholders with actual values from `values.yaml`

```
$ helm template .
---
# Source: Web Application/templates/web-app-pod-template.yaml
apiVersion: v1
kind: Pod
metadata:
  name: web-app
spec:
  containers:
  - image: bmuschko/web-app:1.0.1
    name: web-app
    env:
      - name: DB_HOST
        value: mysql-service
      - name: DB_USER
        value: root
      - name: DB_PASSWORD
        value: password
  ...
```

Installing a Chart from Local Files

Helpful for trying out the deployment before publishing the chart file

```
$ helm install web-app .
```

```
NAME: web-app
```

```
LAST DEPLOYED: Wed Apr 19 11:58:34 2023
```

```
NAMESPACE: default
```

```
STATUS: deployed
```

```
REVISION: 1
```

```
TEST SUITE: None
```

```
$ kubectl get pods,services
```

NAME	READY	STATUS	RESTARTS	AGE
pod/web-app	1/1	Running	0	3m24s

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
service/web-app-service	NodePort	10.109.44.239	<none>	3000:30456/TCP	3m24s

Overriding Values

End user can tweak runtime behavior to personal requirements

```
$ helm install --namespace custom --create-namespace --set service_port=9090 web-app .
```

```
NAME: web-app
```

```
LAST DEPLOYED: Wed Apr 19 11:58:34 2023
```

```
NAMESPACE: custom STATUS: deployed
```

```
REVISION: 1
```

```
TEST SUITE: None
```

Providing a custom
namespace (defaults to
default namespace)

Override values in
values.yaml

Building the Template Files into a Chart Archive File

Archive file is usually published to a chart repository for consumption

```
$ helm package .  
Successfully packaged chart and saved it to:  
/Users/bmuschko/helm/web-app-2.5.4.tgz
```

```
$ ls  
web-app-2.5.4.tgz
```

Lab 10

Implementing, packaging, and installing a custom Helm chart



Application Observability & Maintenance





Probes

Readiness, liveness, and startup probes


What is a Probe?

Detection and potential correction of application runtime issues

- You can configure a container with *health probes* to execute a periodic mini-process that checks for certain conditions.
- Verification methods include running a custom command, sending a HTTP GET request, or opening a TCP socket connection. Pick the method most conducive to your application.
- The verification method defined by a probe is executed by the kubelet.


Readiness Probe

"Is application ready to serve requests?"




Liveness Probe

"Does the application still function without errors?"




Startup Probe

"Legacy application may need longer to start. Hold off on liveness probing."



Are you started? **No**  *Start liveness probe*



Are you started? **Yes**  *Start liveness probe*

Health Verification Methods

Any verification method can be used by any of the probe types

Method	Option	Description
Custom Command	<code>exec.command</code>	Executes a command inside of the container e.g. a cat command and checks its exit code. Kubernetes considers an zero exit code to be successful. A non-zero exit code indicates an error.
HTTP GET Request	<code>httpGet</code>	Sends an HTTP GET request to an endpoint exposed by the application. A HTTP response code in the range of 200 and 399 indicates success. Any other response code is regarded as an error.
TCP socket connection	<code>tcpSocket</code>	Tries to open a TCP socket connection to a port. If the connection could be established, the probing attempt was successful. The inability to connect is accounted for as an error.
gRPC	<code>grpc</code>	The application implements the gRPC Health Checking Protocol . Verifies whether the server is able to handle rpcs.

Health Check Attributes

Any verification method can be used by any of the probe types

Attribute	Default Value	Description
<code>initialDelaySeconds</code>	0	Delay in seconds until first check is executed.
<code>periodSeconds</code>	10	Interval for executing a check (e.g., every 20 seconds).
<code>timeoutSeconds</code>	1	Maximum number of seconds until the check operation times out.
<code>successThreshold</code>	1	Number of successful check attempts until probe is considered successful after a failure.
<code>failureThreshold</code>	3	Number of failures for check attempts before probe is marked failed and takes action.
<code>terminationGracePeriod Seconds</code>	30	Grace period before forcing a container to stop upon failure.

Defining a Readiness Probe

HTTP probes are very helpful for web applications

```
apiVersion: v1
kind: Pod
metadata:
  name: web-app
spec:
  containers:
  - name: web-app
    image: eshop:4.6.3
    readinessProbe:
      httpGet:
        path: /
        port: 8080
      initialDelaySeconds: 5
      periodSeconds: 2
```

Successful if HTTP status
code is between 200 and 399



Defining a Liveness Probe

Querying an event log with a custom command

```
apiVersion: v1
kind: Pod
metadata:
  name: web-app
spec:
  containers:
  - name: web-app
    image: eshop:4.6.3
    livenessProbe:
      exec:
        command:
        - test `find /tmp/heartbeat.txt -mmin -1`
      initialDelaySeconds: 10
      periodSeconds: 5
```

Does the file get updated
periodically by the application
process?



Defining a Startup Probe

Opening a TCP socket connection exposed by the application

```
apiVersion: v1
kind: Pod
metadata:
  name: startup-pod
spec:
  containers:
  - image: httpd:2.4.46
    name: http-server
    startupProbe:
      tcpSocket:
        port: 80
        initialDelaySeconds: 3
        periodSeconds: 15
    livenessProbe:
      ...
```

← Tries to open a TCP socket connection to a port


Using a Named Port

Referencing a port by assigned name in multiple probes

```
ports:
- name: http-port
  containerPort: 8080

readinessProbe:
  httpGet:
    path: /healthz
    port: http-port

livenessProbe:
  httpGet:
    path: /healthz
    port: http-port
```

A diagram consisting of three red rectangular boxes. The first box is around the 'http-port' text in the 'ports' section. The second box is around the 'http-port' text in the 'readinessProbe' section. The third box is around the 'http-port' text in the 'livenessProbe' section. A black line starts from the right side of the first box, goes up and then right to the right side of the second box. From the right side of the second box, a line goes down and then right to the right side of the third box. This illustrates that the named port 'http-port' is referenced by multiple different probes.

Lab 11

Defining probes for a container






Metrics Server

Retrieving, inspecting, and using resource metrics

What is the Metrics Server?

Cluster-wide resource metrics aggregator

- Gathers resource consumption metrics on nodes, e.g. CPU and memory. The kubelet sends those metrics to the Kubernetes API server and exposes them through the [Metrics API](#).
- The metrics server stores data in memory and does not persist data over time.



Installing the Metrics Server

A Kubernetes cluster does not come with the metrics server by default



```
$ minikube addons enable metrics-server  
The 'metrics-server' addon is enabled
```



```
$ kubectl apply -f  
https://github.com/kubernetes-sigs/metrics-server/r  
eleases/latest/download/components.yaml
```

Querying for Metrics

Metrics for nodes and Pods can be rendered with the `top` command

```
$ kubectl top nodes
```

NAME	CPU (cores)	CPU%	MEMORY (bytes)	MEMORY%
controlplane	193m	9%	1239Mi	32%
node01	119m	5%	889Mi	23%

```
$ kubectl top pods
```

NAME	CPU (cores)	MEMORY (bytes)
frontend	38m	85Mi
backend	29m	59Mi

Kubernetes Features that Tap into Metrics

Where else does the metrics server come into play? ([docs](#))

- A Horizontal Pod Autoscaler (HPA) can only determine if Deployment replicas need to be scaled up or down based on available resource consumption determined by metrics server. If the current resource consumption cannot be determined, it is indicated by `<unknown>` in the `TARGETS` column of an HPA.
- Pod scheduling on cluster nodes is based on the definition of resource requests. The scheduler looks at the resources capacity available on a cluster node and decides if the workload can be scheduled for the Pod's resource request.



Troubleshooting Pods and Containers

Techniques for identifying the root cause of a failure

Retrieving High-Level Pod Information

Check the status of a Pod first

```
$ kubectl get pods -n t67
```

NAME	READY	STATUS	RESTARTS	AGE
misbehaving-pod	0/1	ErrImagePull	0	2s

```
$ kubectl get pods -n k31
```

NAME	READY	STATUS	RESTARTS	AGE
successful-pod	1/1	Running	0	5s

Failed to pull
container image

Looks successful on
the surface-level

Common Pod Error Statuses

Poorly-documented in Kubernetes docs, encountered often in practice

Status	Root Cause	Potential Fix
<code>ImagePullBackOff</code> or <code>ErrImagePull</code>	Image could not be pulled from registry.	Check correct image name, check that image name exists in registry, verify network access from node to registry, ensure proper authentication.
<code>CrashLoopBackOff</code>	Application or command run in container crashes.	Check command executed in container, ensure that image can properly execute (e.g., by creating a container with Docker Engine).
<code>CreateContainerConfigError</code>	ConfigMap or Secret referenced by container cannot be found.	Check correct name of the configuration object, verify the existence of the configuration object in the namespace.

Inspecting Events of a Specific Pod

Detailed event information can be found in `Events` section

```
$ kubectl describe pod myapp
...
Events:
  Type          Reason          Age          From          Message
  ----          -
  Normal        Scheduled        <unknown>    default-scheduler    Successfully
  assigned default/secret-pod to minikube
  Warning       FailedMount      3m15s        kubelet, minikube    Unable to
  attach or mount volumes: unmounted volumes=[mysecret], unattached
  volumes=[default-token-bf8rh mysecret]: timed out waiting for the condition
  Warning       FailedMount      68s (x10 over 5m18s)    kubelet, minikube    MountVolume.Setup failed for volume "mysecret" : secret "mysecret" not found
  ...
```


Inspecting Events Across all Pods

Lists events for all Pods in the given namespace

```
$ kubectl get events
```

LAST SEEN	TYPE	REASON	OBJECT	MESSAGE
3m14s	Warning	BackOff	pod/custom-cmd	Back-off restarting failed container ↵
2s	Warning	FailedNeedsStart	cronjob/google-ping	Cannot determine if job needs to be started: too many missed start time (> 100). Set or decrease .spec.startingDeadlineSeconds or check clock skew ↵

Inspecting Container Logs

Application failure may not necessarily surface on the Kubernetes-level

```
$ kubectl create -f crash-loop-backoff.yaml
pod/incorrect-cmd-pod created
```

```
$ kubectl get pods incorrect-cmd-pod
```

NAME	READY	STATUS	RESTARTS	AGE
incorrect-cmd-pod	0/1	CrashLoopBackOff	5	3m20s

```
$ kubectl logs incorrect-cmd-pod
/bin/sh: unknown: not found
```

Opening a Interactive Shell to a Container

Allows for exploring container file system and processes

```
$ kubectl logs failing-pod
```

```
/bin/sh: can't create /root/tmp/curr-date.txt: nonexistent directory
```

```
$ kubectl exec failing-pod -it -- /bin/sh
```

```
# mkdir -p ~/tmp
```

```
# cd ~/tmp
```

```
# ls -l
```

```
total 4
```

```
-rw-r--r-- 1 root root 112 May 9 23:52 curr-date.txt
```

```
# exit
```

Distroless Containers

Some container images do not provide a shell

```
$ kubectl run minimal-pod --image=k8s.gcr.io/pause:3.1
pod/minimal-pod created
```

```
$ kubectl exec minimal-pod -it -- /bin/sh
```

```
OCI runtime exec failed: exec failed: container_linux.go:349:
←
```

```
Starting container process caused "exec: \"/bin/sh\": stat
/bin/sh: no such file or directory": unknown command
terminated with exit code 126
```

Debugging Using an Ephemeral Container

Using a disposable container for troubleshooting distroless containers

```
$ kubectl debug -it minimal-pod --image=busybox
```

```
Defaulting debug container name to debugger-jf98g.
```

```
If you don't see a command prompt, try pressing enter.
```

```
/ # pwd
```

```
/
```

```
/ # exit
```

```
Session ended, resume using 'kubectl attach minimal-pod -c debugger-jf98g -i -t' command when the pod is running ↵
```

You can create a copy of a container in a new Pod with `--copy-to` for more complex debugging scenarios.

Copying a File from the Container

The `tar` binary needs to exist in container

```
$ kubectl exec -it nginx -n code-red -- /bin/sh
# tar cvf error-log.tar /var/log/nginx/error.log
# exit

$ kubectl cp code-red/nginx:error-log.tar error-log.tar

$ ls
error-log.tar
```

Namespace

Pod

File to be copied

Lab 12



Troubleshooting a Pod



Application Environment, Configuration & Security





Custom Resource Definitions (CRDs)

Extending the Kubernetes API by custom primitives


What is a Custom Resource Definition (CRD)?

Extension point for introducing custom API primitives

- Some use cases with custom requirements cannot be fulfilled with the built-in Kubernetes primitives and functionality.
- A CRD allows you to define, create, and persist objects for one or many custom primitives and expose them with the Kubernetes API.
- CRDs are combined with *controllers* to make them actually useful. Controllers interact with the API server and implement the custom logic. The combination of CRDs and controllers is called the *Operator pattern*.

The Operator Pattern

Consists of CRDs and Controllers



Example CRD

Smoke testing against a Service after deployment

- *Requirement:* An web-based application stack is deployed via a Deployment with one or more than one replica. The Service object routes network traffic to the Pods.
- *Desired functionality:* After the deployment happened, we want to run a quick smoke test against the Service's DNS name to ensure that application is operational. The result of the smoke test will be sent to an external service for the purpose of rendering it on the UI dashboard.
- *Goal:* Implement the Operator pattern (CRD and controller) for the use case.

CRD YAML Manifest

Defines the schema for the custom object

```
apiVersion: apiextensions.k8s.io/v1
kind: CustomResourceDefinition
metadata:
  name: smoketests.stable.bmuschko.com
spec:
  group: stable.bmuschko.com
  versions:
    - name: v1
      served: true
      storage: true
      schema:
        openAPIV3Schema:
          type: object
          properties:
            spec:
              type: object
              properties:
                service:
                  type: string
                ...
  scope: Namespaced
  names:
    plural: smoketests
    singular: smoketest
    kind: SmokeTest
    shortNames:
      - st
```

← Combination of <plural>.<group>

← Versions supported by CRD

← Attributes for custom type

← Identifiers for custom type

Custom Resource Object YAML Manifest

Instantiation of CRD kind

```
apiVersion:
stable.bmuschko.com/v1
kind: SmokeTest
metadata:
  name: backend-smoke-test
spec:
  service: backend
  path: /health
  timeout: 600
  retries: 3
```

← Group and Version of custom kind

← Kind defined by CRD

← Attributes and values that make
custom kind configurable

Creating the CRD and Custom Object

CRD object needs to be created before custom objects

```
$ kubectl apply -f loadtest-resource.yaml
customresourcedefinition.apiextensions.k8s.io/smoketests.stable.bmuschko.com created
```

```
$ kubectl apply -f loadtest.yaml
smoketest.stable.bmuschko.com/backend-smoke-test created
```

Discovering CRDs

Once installed, you can list all available CRDs

```
$ kubectl api-resources --api-group=stable.bmuschko.com
```

NAME	SHORTNAMES	APIVERSION	NAMESPACED	KIND
smoketests	st	stable.bmuschko.com/v1	true	SmokeTest

```
$ kubectl get crds
```

NAME	CREATED AT
smoketests.stable.bmuschko.com	2023-05-04T14:49:40Z

Controller Implementation

The actual logic for smoke tests lives in the controller

- You can interact with the custom resource object using CRUD (create/read/update/delete) operations; however, no smoke test logic will actually be initiated.
- A controller acts as a reconciliation process by inspecting the state through the API server. At runtime, it polls for new `SmokeTest` objects and sends the result to the external service.
- Controllers can use one of the Kubernetes [client libraries](#), written in Go or Python, to access custom resources.

Lab 13

Defining and interacting with a CRD






Authentication and Authorization

Access control for the Kubernetes API

Processing a Request to the API Server

Any client request to the API server needs to pass four phases



API Request Processing Phases

Every request has to pass those phases in the following order

- *Authentication*: Validates the identity of the caller using a supported authentication strategy, e.g. client certificates or bearer tokens.
- *Authorization*: Determines if the identity provided in the first stage can access the verb and HTTP path request. This is usually implemented with RBAC.
- *Admission Controller*: Verifies if the request is well-formed and potentially needs to be modified before the request is processed. Ensures that the resource included in the request is valid (could also be implemented as part of admission control).

Authentication via Credentials in Kubeconfig

The kubeconfig file at `$HOME/.kube/config` is evaluated by `kubectl`

```
apiVersion: v1
kind: Config
clusters:
- cluster:
    certificate-authority: /Users/bmuschko/.minikube/ca.crt
    extensions:
    - extension:
        last-update: Thu, 04 May 2023 08:48:09 MDT
        provider: minikube.sigs.k8s.io
        version: v1.30.1
        name: cluster_info
    server: https://127.0.0.1:58936
  name: minikube
contexts:
- context:
    cluster: minikube
    user: bmuschko
    name: bmuschko
  current-context: bmuschko
users:
- name: bmuschko
  user:
    client-certificate-data: REDACTED
    client-key-data: REDACTED
```

API server endpoint for
connecting with cluster

Groups access parameters
under a convenient name

Client certificate for user

Managing Kubeconfig and Current Context

Interaction with `kubeconfig` configuration

```
$ kubectl config view
```

```
apiVersion: v1  kind: Config  
current-context: bmuschko  
...
```

← Renders the contents of
the kubeconfig file

```
$ kubectl config current-context
```

```
bmuschko
```

← Shows the
currently-selected context

```
$ kubectl config use-context johndoe
```

```
Switched to context "johndoe".
```

← Switches to the context
defined in the kubeconfig

```
$ kubectl config current-context
```

```
johndoe
```


What is RBAC?

Restricting access control for clients interacting with API server

- Defines policies for users, groups, and processes by allowing or disallowing access to manage API resources.
- Enabling and configuring RBAC is mandatory for any organization with a strong emphasis on security.
- *Example:* “The human user John is only allowed to list and create Pods and Deployments, but nothing else.”


RBAC High-Level Overview

Three key elements for understanding concept



Involved RBAC Primitives

Restrict access to certain operations of API resources for subjects



Namespace and Cluster-Wide Permissions

Defining permission scopes

- The *Role* maps API resources to verbs for a single namespace.
- The *RoleBinding* maps a reference of a Role to one or many subjects for a single namespace.
- To apply permissions across all namespaces of a cluster or cluster-wide resources, use the corresponding API primitives *ClusterRole* and *ClusterRoleBinding*. The configuration options of the manifests are the same.

Creating a Role with Imperative Approach

Defines verbs and resources in comma-separated list

```
$ kubectl create role read-only --verb=list,get,watch --resource=pods,deployments,services  
role.rbac.authorization.k8s.io/read-only created
```

Resources: Primitives the operations should apply to

Operations: Only allow listing, getting, watching

Role YAML Manifest

Can define a list of rules in an array

```
apiVersion: rbac.authorization.k8s.io/v1
kind: Role
metadata:
  name: read-only
rules:
- apiGroups: [""]
  resources: ["pods", "services"]
  verbs: ["list", "get", "watch"]
- apiGroups: ["apps"]
  resources: ["deployments"]
  verbs: ["list", "get", "watch"]
```

Resources with groups need to be spelled out explicitly (in this case `apps/Deployment`)

Getting Role Details

Maps objects of a Kubernetes primitive to verbs

```
$ kubectl describe role read-only
```

```
Name:          read-only
```

```
Labels:        <none>
```

```
Annotations:   <none>
```

```
PolicyRule:
```

Resources	Non-Resource URLs	Resource Names	Verbs
-----	-----	-----	-----
pods	[]	[]	[list get watch]
services	[]	[]	[list get watch]
deployments.apps	[]	[]	[list get watch]

Creating a RoleBinding with Imperative Approach

Maps subject to Role

```
$ kubectl create rolebinding read-only-binding --role=read-only --user=johndoe  
rolebinding.rbac.authorization.k8s.io/read-only-binding created
```

Subject: Binds a user to
the Role

Role: Reference the name of
the object

RoleBinding YAML Manifest

Roles can be mapped to multiple subjects if needed

```
apiVersion: rbac.authorization.k8s.io/v1
kind: RoleBinding
metadata:
  name: read-only-binding
roleRef:
  apiGroup: rbac.authorization.k8s.io
  kind: Role
  name: read-only
subjects:
- apiGroup: rbac.authorization.k8s.io
  kind: User
  name: johndoe
```

← Reference to Role

← Reference to User

Getting RoleBinding Details

Only shows mapping between Role and subjects

```
$ kubectl describe rolebinding read-only-binding
```

```
Name:          read-only-binding
```

```
Labels: <none>
```

```
Annotations: <none>
```

```
Role:
```

```
  Kind:      Role
```

```
  Name:      read-only
```

```
  Subjects:
```

Kind	Name	Namespace
------	------	-----------

-----	-----	
-------	-------	--

User	johndoe	
------	---------	--

What is a Service Account?

Non-human request to Kubernetes API from a process

- Processes running outside of Kubernetes or processes running inside of a container may need to interact with the Kubernetes API. A Service Account allows for authenticating with the API server through an authentication token.
- Authorization is controlled through RBAC and assigned to the Service Account.
- If not assigned explicitly, a Pod uses the `default` Service Account. The `default` Service Account has the same permissions as an unauthenticated user.
- *Example:* “I have a CI/CD process that retrieves Pod information by making calls to the Kubernetes API.”

Service Account as Subject

Pod assign the Service Account by name



Using the kubernetes Service

Convenient way to get API server endpoint from within the cluster

```
$ kubectl get service kubernetes
```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)
kubernetes	ClusterIP	10.96.0.1	AGE	
			<none>	443/TCP
			13h	

```
$ kubectl run api-call --image=alpine:3.16.2 -it --rm --restart=Never -- https://kubernetes.default.svc.cluster.local/api/v1/namespaces/k97/pods
```


API endpoint for
listing all Pods in
namespace k97

Service Account YAML Manifest

In a Pod manifest, refer to the Service Account by name

```
apiVersion: v1
kind: ServiceAccount
metadata:
  name: sa-api
  namespace: k97
```

```
apiVersion: v1
kind: Pod
metadata:
  name: list-pods
  namespace: k97
spec:
  serviceAccountName: sa-api
  ...
```



Accessing Service Account Authentication Token

Automounted at specific Volume mount path

```
$ kubectl describe pod list-pods -n k97
...
Containers:
  pods:
    ...
    Mounts:
      /var/run/secrets/kubernetes.io/serviceaccount from kube-api-access-xnkwk (ro)
...

$ kubectl exec -it list-pods -n k97 -- /bin/sh
# cat /var/run/secrets/kubernetes.io/serviceaccount/token
eyJhbGciOiJSUzI1Ni...
# exit
```

Using Service Account in Pod

Enables authentication token auto-mounting by default

```
apiVersion: v1
kind: Pod
metadata:
  name: list-pods
  namespace: k97
spec:
  serviceAccountName: sa-api
  containers:
  - name: pods
    image: alpine/curl:3.14
    command: ['sh', '-c', 'while true; do curl -s -k -m 5 -H
      "Authorization: Bearer $(cat /var/run/secrets/kubernetes.io/
      serviceaccount/token)" https://kubernetes.default.svc.cluster.
      local/api/v1/namespaces/k97/pods; sleep 10; done']
```

Assigned name of
Service Account


List all Pods in the
namespace k97
via an API call

RoleBinding YAML Manifest

Default RBAC policies don't span beyond kube-system

```
apiVersion: rbac.authorization.k8s.io/v1
kind: RoleBinding
metadata:
  name: serviceaccount-pod-rolebinding
  namespace: k97
roleRef:
  apiGroup: rbac.authorization.k8s.io
  kind: ClusterRole
  name: list-pods-clusterrole
subjects:
- apiGroup: rbac.authorization.k8s.io
  kind: ServiceAccount
  name: sa-api
```

Maps the Service
Account as a
subject



Lab 14

Regulating Access to API Resources with RBAC



Resource Management

Resource requests and limits, resource quotas, limit ranges

Managing Resources for Objects and Namespaces

It's best practice to make a statement about resource consumption

- Containers can define a minimum amount of resources needed to run the application, as well as the maximum amount of resources the application is allowed to consume.
- Application developers should determine the right-sizing with load tests or at runtime by monitoring the resource consumption.
- Enforcement of resources can be constrained on a namespace-level with a *ResourceQuota*.
- A *LimitRange* is a policy that constrains or defaults the resource allocations for a single object (e.g., for a Pod or PersistentVolumeClaim).

Resource Units in Kubernetes

CPU units and memory as fixed-point number or power-of-two equivalents

- Kubernetes measures CPU resources in millicores and memory resources in bytes. That's why you might see resources defined as `600m` or `100Mib`.
- For a deep dive on those resource units, it's worth cross-referencing the section ["Resource units in Kubernetes"](#) in the official documentation.


Container Resource Requests

Definition and runtime effects

- Containers can define a minimum amount of resources needed to run the application via `spec.containers[].resources.requests`.
- Resource types include CPU, memory, huge page, ephemeral storage.
- If Pod cannot be scheduled on any node due to insufficient resource availability, you may see a `PodExceedsFreeCPU` or `PodExceedsFreeMemory` status.

Example Scheduling Scenario

Node's resource capacity needs to be able to fulfill it



Container Resource Requests Options

Configuration options and example values

YAML Attribute	Description	Example Value
<code>spec.containers[].resources.requests.cpu</code>	CPU resource type	500m (five hundred millicpu)
<code>spec.containers[].resources.requests.memory</code>	Memory resource type	64Mi (2 ²⁶ bytes)
<code>spec.containers[].resources.requests.hugepages-<size></code>	Huge page resource type	hugepages-2Mi: 60Mi
<code>spec.containers[].resources.requests.ephemeral-storage</code>	Ephemeral storage resource type	4Gi

Container Resource Requests YAML Manifest

Minimum amount of resources defined for a container

```
apiVersion: v1
kind: Pod
metadata:
  name: rate-limiter
spec:
  containers:
  - name: business-app
    image: bmuschko/nodejs-business-app:1.0.0
    ports:
    - containerPort: 8080
    resources:
      requests:
        memory: "256Mi"
        cpu: "1"
```


Container Resource Limits

Definition and runtime effects

- Containers can define a maximum amount of resources allowed to consumed by the application via `spec.containers[].resources.limits`.
- Resource types include CPU, memory, huge page, ephemeral storage.
- Container runtime decides how to handle situation where application exceeds allocated capacity, e.g. termination of application process.

Container Resource Limits Options

Configuration options and example values

YAML Attribute	Description	Example Value
<code>spec.containers[].resources.limits.cpu</code>	CPU resource type	500m (five hundred millicpu)
<code>spec.containers[].resources.limits.memory</code>	Memory resource type	64Mi (2 ²⁶ bytes)
<code>spec.containers[].resources.limits.hugepages-<size></code>	Huge page resource type	hugepages-2Mi: 60Mi
<code>spec.containers[].resources.limits.ephemeral-storage</code>	Ephemeral storage resource type	4Gi

Container Resource Limits YAML Manifest

Do not allow more than the allotted resource amounts

```
apiVersion: v1
kind: Pod
metadata:
  name: rate-limiter
spec:
  containers:
  - name: business-app
    image: bmuschko/nodejs-business-app:1.0.0
    ports:
    - containerPort: 8080
    resources:
      limits:
        memory: "512Mi"
        cpu: "2"
```

Pod Scheduling Details

After scheduling a Pod, you can check on runtime information

```
$ kubectl get pod rate-limiter -o yaml | grep nodeName:
  nodeName: worker-3
```

```
$ kubectl describe node worker-3
```

```
...
```

Non-terminated Pods:	(3 in total)			
Namespace	Name	CPU Requests	CPU Limits	...
-----	----	-----	-----	...
default	rate-limiter	1250m (62%)	0 (0%)	...

Container Resource Requests/Limits YAML Manifest

It is considered best practice to define requests *and* limits for every container

```
spec:
  containers:
  - name: business-app
    resources:
      requests:
        memory: "256Mi"
        cpu: "1"
      limits:
        memory: "512Mi"
        cpu: "2"
```

What is a ResourceQuota?

Defines resource constraints per namespace

- Constraints that limit aggregate resource consumption or limit the quantity of objects (e.g., number of Pods or Secrets) that can be created.
- `requests.cpu/requests.memory`: Across all pods in a non-terminal state, the sum of CPU/memory requests cannot exceed this value.
- `limits.cpu/limits.memory`: Across all pods in a non-terminal state, the sum of CPU/memory limits cannot exceed this value.

ResourceQuota YAML Manifest

Limits # of objects and aggregate container resource consumption

```
apiVersion: v1
kind: ResourceQuota
metadata:
  name: awesome-quota
  namespace: team-awesome
spec:
  hard:
    pods: 2
    requests.cpu: "1"
    requests.memory: 1024Mi
    limits.cpu: "4"
    limits.memory: 4096Mi
```

No Requests/Limits Definition

Scheduler rejects the creation of the object

```
$ kubectl create -f nginx-pod.yaml -n team-awesome
```

```
Error from server (Forbidden): error when creating "nginx-pod.yaml": ␣  
pods "nginx" is forbidden: failed quota: awesome-quota: must specify ␣  
limits.cpu,limits.memory,requests.cpu,requests.memory
```


Exceeds Requests/Limits Definition

Scheduler rejects the creation of the object

```
$ kubectl create -f nginx-pod.yaml -n team-awesome
```

```
Error from server (Forbidden): error when creating "nginx-pod.yaml": ␣  
pods "nginx" is forbidden: exceeded quota: awesome-quota, requested: ␣  
pods=1,requests.cpu=500m,requests.memory=512Mi, used: pods=2, ␣  
requests.cpu=1,requests.memory=1024Mi, limited: pods=2, ␣  
requests.cpu=1,requests.memory=1024Mi
```

Inspecting a ResourceQuota

Renders consumed resources and defined hard limits

```
$ kubectl describe resourcequota awesome-quota -n team-awesome
```

```
Name:                awesome-quota
```

```
Namespace:           team-awesome
```

```
Resource             Used      Hard
```

```
-----
```

```
----
```

```
----
```

```
limits.cpu           2         4
```

```
limits.memory        2048m     4096m
```

```
pods                 2         2
```

```
requests.cpu          1         1
```

```
requests.memory      1024m     1024m
```

What is a LimitRange?

Constraints or defaults the resource allocations for an object type

- Enforces minimum and maximum compute resources usage per Pod or container in a namespace.
- Enforces minimum and maximum storage request per PersistentVolumeClaim in a namespace.
- Enforces a ratio between request and limit for a resource in a namespace.
- Set default requests/limits for compute resources in a namespace and automatically injects them into containers at runtime.

LimitRange YAML Manifest

Sets defaults, minimum and maximum CPU allocation for containers

```
apiVersion: v1
kind: LimitRange
metadata:
  name:
cpu-resource-constraint
spec:
  limits:
  - default:
      cpu: 500m
    defaultRequest:
      cpu: 500m
    max:
      cpu: "1"
    min:
      cpu: 100m
  type: Container
```

- ← Defines default limits
- ← Defines default requests
- ← Minimum and maximum resource range

Automatically Uses Container Defaults for Pod

Container doesn't provide any resource requests/limits

```
$ kubectl create -f pod.yaml
pod/rate-limiter created

$ kubectl describe pod rate-limiter
...
Containers:
  business-app:
    ...
    Limits:
      cpu: 500m
    Requests:
      cpu: 500m
  ...
...
```

Rejected Pod Creation for Exceeded Limits

Requests a higher maximum for CPU resources than defined by LimitRange

```
$ kubectl create -f pod.yaml
```

```
Error from server (Forbidden): error when creating "pod.yaml": pods ↵  
"rate-limiter" is forbidden: maximum cpu usage per Container is 1, ↵  
but limit is 2
```

Lab 15

Resources requirements





ConfigMaps and Secrets

Defining and consuming configuration data


Defining Configuration Data

Control runtime behavior with centralized information

- Stored as key-value pairs in dedicated Kubernetes primitives with a lifecycle decoupled from consuming Pod.
- *ConfigMap*: Plain-text values suited for configuring applications e.g. flags, or URLs.
- *Secret*: Base64-encoded values suited for storing sensitive data like passwords, API keys, or SSL certificates. **Values are not encrypted!**
- ConfigMap and Secret objects are stored in etcd in unencrypted form by default. Encryption of data in etcd is [configurable](#).

Consuming Configuration Data

Applicable to ConfigMap and Secret



Creating a ConfigMap with Imperative Approach

Key-value pairs can be parsed from different sources

```
$ kubectl create configmap db-config --from-literal=db=staging
configmap/db-config created
```

```
$ kubectl create configmap db-config --from-env-file=config.env
configmap/db-config created
```

```
$ kubectl create configmap db-config --from-file=config.txt
configmap/db-config created
```

```
$ kubectl create configmap db-config --from-file=app-config
configmap/db-config created
```

Source Options for Data Parsed by a ConfigMap

Configuration options and example values

Option	Description	Example Value
<code>--from-literal</code>	Literal values, which are key-value pairs as plain text.	<code>--from-literal=locale=en_US</code>
<code>--from-env-file</code>	A file that contains key-value pairs and expects them to be environment variables.	<code>--from-env-file=config.env</code>
<code>--from-file</code>	Huge page resource type.	<code>--from-file=app-config.json</code>
<code>--from-file</code>	A directory with one or many files.	<code>--from-file=config-dir</code>

Creating a ConfigMap from Literal Values

Imperative creation by pointing to plain-text key-value pairs

```
$ kubectl create configmap  
  db-config  
  --from-literal=DB_HOST=mysql-service  
  --from-literal=DB_USER=backend  
configmap/db-config created
```

← Parameter can
be repeated
multiple times

ConfigMap YAML Manifest

Key-value pair definition under `data` attribute


```
apiVersion: v1
kind: ConfigMap
metadata:
  name: db-config
data:
  DB_HOST: mysql-service
  DB_USER: backend
```

← Follows typical
naming conventions
for environment
variables


Consuming a ConfigMap as Environment Variables

Controlling runtime behavior with the help of simple values

```
apiVersion: v1
kind: Pod
metadata:
  name: backend
spec:
  containers:
  - image: bmuschko/web-app:1.0.1
    name: backend
    envFrom:
    - configMapRef:
      name: db-config
```



```
$ kubectl exec -it backend -- env
DB_HOST=mysql-service
DB_USER=backend
```



Reassigning Environment Variable Keys

Keys can be renamed or reformatted to make them suitable for consumption

```
apiVersion: v1  kind: Pod
metadata:
  name: backend
spec:
  containers:
    - image: bmuschko/web-app:1.0.1
```

```
name: backend
```

```
env:
```

```
- name: DATABASE_URL
  valueFrom:
    configMapKeyRef:
      name: backend-config
      key: database_url
```

```
$ kubectl exec -it backend -- env
DATABASE_URL=jdbc:postgresql:↵
//localhost/test
```


Creating a ConfigMap from a JSON file

Imperative creation by pointing to file

```
$ kubectl create configmap db-config --from-file=db.json  
configmap/db-config created
```

db.json



```
{  
  "db": {  
    "host": "mysql-service",  
    "user": "backend"  
  }  
}
```

ConfigMap YAML Manifest

Key-value pair definition under `data` attribute


```
apiVersion: v1
kind: ConfigMap
metadata:
  name: db-config
data:
  db.json: |-
    {
      "db": {
        "host": "mysql-service",
        "user": "backend"
      }
    }
```

← The file name
becomes the key, the
contents of the file
become the value

Consuming a ConfigMap as a Volume

A good choice for processing a machine-readable configuration file

```
apiVersion: v1
kind: Pod
metadata:
  name: backend
spec:
  containers:
    - name: backend
      image:
bmuschko/web-app:1.0.1
      volumeMounts:
        - name: db-config-volume
          mountPath: /etc/config
  volumes:
    - name: db-config-volume
      configMap:
        name: db-config
```

A diagram illustrating how a Kubernetes ConfigMap is consumed as a volume. On the left, a Pod configuration is shown. The 'volumeMounts' section of the 'backend' container is highlighted with a red box, showing it mounts 'db-config-volume' to '/etc/config'. The 'volumes' section shows 'db-config-volume' is a 'ConfigMap' named 'db-config'. An arrow points from this configuration to a terminal window on the right. The terminal shows the command '\$ kubectl exec -it backend -- /bin/sh' and the output of 'ls -l /etc/config' showing 'db.json'. A red box highlights the output of 'cat /etc/config/db.json', which is a JSON object with 'db' containing 'host' and 'user' fields.

```
$ kubectl exec -it backend -- /bin/sh
# ls -l /etc/config
db.json
# cat /etc/config/db.json
{
  "db": {
    "host": "mysql-service",
    "user": "backend"
  }
}
```

Consuming Changed ConfigMap Data

Containers will not refresh consumed data upon a change

- The key-value pairs of a ConfigMap can be changed for a live object.
- Changed ConfigMap key-value pairs consumed by containers will not be reloaded automatically for an already-running Pod.
- The application code needs to implement logic to either reload the configuration data periodically or on-demand.

Creating a Secret with Imperative Approach

Parsed values are base64-encoded automatically

```
$ kubectl create secret generic db-creds --from-literal=pwd=s3cre!  
secret/db-creds created
```

```
$ kubectl create secret docker-registry my-secret <  
  --from-file=.dockerconfigjson=~/.docker/config.json  
secret/my-secret created
```

```
$ kubectl create secret tls accounting-secret --cert=accounting.crt <  
  --key=accounting.key  
secret/accounting-secret created
```

Imperative CLI Options for Creating a Secret

You will need to provide one of the options

Option	Description	Type
<code>generic</code>	Creates a secret from a file, directory, or literal value.	Opaque
<code>docker-registry</code>	Creates a secret for use with a Docker registry, e.g. to pull images from a private registry when requested by a Pod.	<code>kubernetes.io/dockercfg</code>
<code>tls</code>	Creates a TLS secret.	<code>kubernetes.io/tls</code>

Source Options for Data Parsed by a Generic Secret

Configuration options and example values

Option	Description	Example Value
<code>--from-literal</code>	Literal values, which are key-value pairs as plain text.	<code>--from-literal=password=secret</code>
<code>--from-env-file</code>	A file that contains key-value pairs and expects them to be environment variables.	<code>--from-env-file=config.env</code>
<code>--from-file</code>	Huge page resource type.	<code>--from-file=id_rsa=~/.ssh/id_rsa</code>
<code>--from-file</code>	A directory with one or many files.	<code>--from-file=config-dir</code>

Specialized Secret Types

Can be set with `--type` CLI option, or the `type` attribute in manifest ([docs](#))

Type	Description
<code>kubernetes.io/basic-auth</code>	Credentials for basic authentication
<code>kubernetes.io/ssh-auth</code>	Credentials for SSH authentication
<code>kubernetes.io/service-account-token</code>	ServiceAccount token
<code>bootstrap.kubernetes.io/token</code>	Node bootstrap token data

Creating a Secret from Literal Values

Imperative command will base64-encode values automatically

```
$ kubectl create secret  
generic  
db-creds  
--from-literal=pwd=s3cre!  
secret/db-creds created
```

← Plain-text value

Secret YAML Manifest

Assigned values need to be provided in base64-encoded form


```
$ echo -n 's3cre!' | base64  
czNjcUUh
```

```
apiVersion: v1  
kind: Secret  
metadata:  
  name: mysecret  
type: Opaque  
data:  
  pwd: czNjcUUh
```

Consuming a Secret as Environment Variables

Accessed values are base64-decoded

```
apiVersion: v1
kind: Pod
metadata:
  name: backend
spec:
  containers:
  - image: bmuschko/web-app:1.0.1
    name: backend
    envFrom:
    - secretRef:
      name: mysecret
```



```
$ kubectl exec -it backend -- env
pwd=s3cre!
```



Creating a Secret from a SSH Private Key File

Imperative creation by pointing to file

```
$ kubectl create secret generic secret-ssh-auth  
  --type=kubernetes.io/ssh-auth --from-file=ssh-privatekey  
secret/secret-ssh-auth created
```

ssh-privatekey



```
-----BEGIN RSA PRIVATE KEY-----  
  
Proc-Type: 4,ENCRYPTED  
DEK-Info:  
AES-128-CBC,8734C9153079F2E8497C8075289EBBF1  
...  
-----END RSA PRIVATE KEY-----
```

Consuming a Secret as a Volume

A good choice for processing a machine-readable configuration file

```
apiVersion: v1
kind: Pod
metadata:
  name: backend
spec:
  containers:
    - name: backend
      image: bmuschko/web-app:1.0.1
      volumeMounts:
        - name: ssh-volume
          mountPath: /var/app
          readOnly: true
      volumes:
        - name: ssh-volume
          secret:
            secretName: secret-ssh-auth
```

```
$ kubectl exec -it backend -- /bin/sh
# ls -l /var/app
ssh-privatekey
# cat /var/app/ssh-privatekey
-----BEGIN RSA PRIVATE KEY-----
Proc-Type: 4,ENCRYPTED
DEK-Info:
AES-128-CBC,8734C9153079F2E8497C8075289E
BBF1
...
-----END RSA PRIVATE KEY-----
```


Plain-Text Secret Values

The `stringData` attribute allows for plain-text values

secret.yaml

```
apiVersion: v1
kind: Secret
metadata:
  name: secret-basic-auth
type: kubernetes.io/basic-auth
stringData:
  username: admin
  password: t0p-Secret
```

kubectl create -f
secret.yaml



Live object

```
apiVersion: v1
kind: Secret
metadata:
  name: secret-basic-auth
type: kubernetes.io/basic-auth
data:
  username: YWRtaW4=
  password: dDBwLVNlY3JldA==
```

Lab 16

Creating and consuming ConfigMap and Secret





Security Context

Privilege and access control for a Pod

What is a Security Context?

Defines security settings for containers of a Pod

- Modeled as a set of attributes within the Pod specification, e.g. for configuring Linux capabilities, file access control, privileges to parent process or host system.
- Can be applied to all containers in Pod with `spec.securityContext` or for individual containers with `spec.containers[].securityContext`.
- Some of the attributes on the Pod- and container-level are the same. Container-level definition of an attribute takes precedence.
- *Example:* “This container needs to run with a non-root user.”

Security Context API

All attributes can be discovered in documentation for the Kubernetes version

API Type	Description
<u>PodSecurityContext</u>	Defines Pod-level security attributes.
<u>SecurityContext</u>	Defines container-level security attributes.

Defining a Security Context

Pod- and/or container-level definition

```
apiVersion: v1
kind: Pod
metadata:
  name: secured-pod
spec:
  securityContext:
    fsGroup: 3000
  volumes:
  - name: data-vol
    emptyDir: {}
  containers:
  - image: nginx:1.18.0
    name: secured-container
    volumeMounts:
    - name: data-vol
      mountPath: /data/app
```

Defined on the Pod-level (applies to all containers)

Inspecting the Runtime Behavior

Rendering the file system Linux group ID

```
$ kubectl exec secured-pod -it -- /bin/sh
# cd /data/app
# touch hello.txt
# ls -l
total 0
-rw-r--r-- 1 root 3000 0 Apr 30 18:27 hello.txt
# exit
```



File system group ID is
assigned automatically

Overriding a Value on the Container-Level

Container-level definition wins

```
apiVersion: v1
kind: Pod
metadata:
  name: non-root-success
spec:
  securityContext:
    runAsNonRoot: false
  containers:
    - image: bitnami/nginx:1.23.1
      name: nginx
      securityContext:
        runAsNonRoot: true
```

← Takes precedence

Inspecting the Runtime Behavior

Rendering the current user ID and group ID

```
$ kubectl exec non-root-success -it -- /bin/sh
$ id
uid=1001 gid=0(root) groups=0(root)
$ exit
```



Non-root user ID

Lab 17

Defining a Security Context for a Pod



Services & Networking





Services

Providing a stable network endpoint to Pods


What is a Service?

Stable endpoint for routing traffic to a set of Pods

- A Pod's virtual IP address is not considered stable over time. A Pod restart leases a new IP address.
- Building a microservices architecture on Kubernetes requires network endpoints that are stable over time.
- A Service provides discoverable names and load balancing to a set of Pods.


Request Routing

"How does a Service decide which Pod to forward the request to?"



Port Mapping

"How to map the Service port to the container port in a Pod?"



Creating a Service with Imperative Approach

Needs to provide the service type and ports

```
$ kubectl run echoserver␣  
  --image=k8s.gcr.io/echoserver:1.10␣  
  --port=8080␣  
pod/echoserver created
```

Exposed container port

```
$ kubectl create␣  
  service clusterip␣  
  echoserver␣  
  --tcp=80:8080␣  
service/echoserver created
```

Incoming and outgoing port

```
$ kubectl run echoserver␣  
  --image=k8s.gcr.io/echoserver:1.10␣  
  --port=8080␣  
  --expose␣  
service/echoserver created  
pod/echoserver created
```

Shortcut for creating a Pod + Service

Service YAML Manifest

Needs to provide the service type and ports

```
apiVersion: v1
kind: Service
metadata:
  name: echoserver
spec:
  selector:
    app: echoserver
  ports:
  - port: 80
    targetPort: 8080
    protocol: TCP
```

Label selection for Pods

Mapping of incoming to outgoing port


Service Types

The following table shows a limited list of Service types

Type	Description
ClusterIP	Exposes the service on a cluster-internal IP. Only reachable from Pods within the cluster. Default if value for <code>spec.type</code> has not been assigned in manifest.
NodePort	Exposes the service on each node's IP at a static port. Accessible from outside of the cluster.
LoadBalancer	Exposes the service externally using a cloud provider's load balancer.


Service Type Inheritance

Services build on top of each other



ClusterIP Service Type

Only reachable from within the cluster, e.g. another Pod



ClusterIP Service YAML Manifest at Runtime

ClusterIP service type has been assigned if not already provided

```
apiVersion: v1
kind: Service
metadata:
  name: echoserver
spec:
  type: ClusterIP
  clusterIP: 10.96.254.0
  selector:
    app: echoserver
  ports:
    - port: 80
      targetPort: 8080
      protocol: TCP
```

Accessing a ClusterIP Service

Use the combination of cluster IP and incoming service port

```
$ kubectl get pod,service
```

NAME	READY	STATUS	RESTARTS	AGE
pod/echoserver	1/1	Running	0	23s

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
service/echoserver	ClusterIP	10.96.254.0	<none>	8080/TCP	8s

```
$ kubectl run tmp --image=busybox --restart=Never -it --rm -- wget 10.96.254.0:8080
Connecting to 10.96.254.0:8080 (10.96.254.0:8080)
...
```

The diagram illustrates the connection from the `wget` command in the terminal to the service's cluster IP and port. A red box highlights the IP `10.96.254.0` in the `CLUSTER-IP` column of the service table. Another red box highlights the port `8080` in the `PORT(S)` column. A third red box highlights the `10.96.254.0:8080` in the `wget` command. Arrows show the flow of information: one arrow points from the `CLUSTER-IP` box to the IP part of the `wget` box, and another arrow points from the `PORT(S)` box to the `:8080` part of the `wget` box.

Discovering the Service by Environment Variables

Application code running in container can get access to Service endpoint

```
$ kubectl exec -it echoserver -- env  
ECHOSERVER_SERVICE_HOST=10.96.254.0  
ECHOSERVER_SERVICE_PORT=8080  
...
```

Discovering the Service by DNS Lookup

Kubernetes' DNS service maps cluster IP address to Service name

```
$ kubectl run tmp --image=busybox --restart=Never -it --rm ↵  
-- wget echoserver.default.svc.cluster.local:8080  
Connecting to echoserver.default.svc.cluster.local:8080 (10.96.254.0:8080)  
...
```

Service name


Namespace

Resource type

Cluster-specific suffix

NodePort Service Type

Can be resolved from outside of the Kubernetes cluster



NodePort Service YAML Manifest at Runtime

Service type is NodePort, node port has been assigned

```
apiVersion: v1
kind: Service
metadata:
  name: echoserver
spec:
  type: NodePort
  clusterIP: 10.96.254.0
  selector:
    app: echoserver
  ports:
    - port: 80
      targetPort: 8080
      nodePort: 30158
      protocol: TCP
```

Accessing a NodePort Service

Use the combination of node IP and statically-assigned port

```
$ kubectl get pod,service
```

NAME	READY	STATUS	RESTARTS	AGE
pod/echoserver	1/1	Running	0	23s

NAME	AGE	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)
service/echoserver		NodePort	10.96.254.0	<none>	8080:30158/TCP 8s

```
$ kubectl get nodes -o ←
```

```
jsonpath='{ $.items[*].status.addresses[?(@.type=="InternalIP")].address }'
```

```
192.168.64.15
```


```
$ wget 192.168.64.15:30158
```

```
Connecting to 192.168.64.15:30158... connected.
```

```
...
```


LoadBalancer Service Type

Routes traffic from existing, external load balancer to Service



LoadBalancer Service YAML Manifest at Runtime

Service type is LoadBalancer, external IP address has been assigned

```
apiVersion: v1
kind: Service
metadata:
  name: echoserver
spec:
  type: LoadBalancer
  clusterIP: 10.96.254.0
  loadBalancerIP: 10.109.76.157
  selector:
    app: echoserver
  ports:
    - port: 80
      targetPort: 8080
      nodePort: 30158
      protocol: TCP
```

Accessing a LoadBalancer Service


Use the combination of external IP and statically-assigned port

```
$ kubectl get pod,service
```

NAME	READY	STATUS	RESTARTS	AGE
pod/echoserver	1/1	Running	0	23s

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)
service/echoserver	LoadBalancer	10.109.76.157	10.109.76.157	8080:30158/TCP

```
$ wget 10.109.76.157:8080
Connecting to 10.109.76.157:8080... connected.
...
```

A diagram consisting of two arrows. The first arrow starts from the 'EXTERNAL-IP' field (10.109.76.157) in the service table and points down to the 'wget' command. The second arrow starts from the 'PORT(S)' field (8080:30158/TCP) in the service table and points down to the same 'wget' command, specifically to the port number 8080.

Lab 18

Routing traffic to Pods
from inside and outside
of a cluster





Ingress

Providing external access to Services via HTTP(S)

What is an Ingress?

Route HTTP(S) traffic from the outside of the cluster to Service(s)

- It's not a specific Service type, nor should it be confused with the Service type `LoadBalancer`.
- Defines rules for mapping a URL context path to one or many Service objects.
- [TLS termination](#) (support for HTTPS) needs to be configured explicitly. By default, only HTTP is supported.


What is an Ingress Controller?

Evaluating Ingress rules

- An Ingress cannot work without an Ingress controller. The Ingress controller evaluates the collection of rules defined by an Ingress that determine traffic routing.
- Multiple Ingress controller can be deployed to a cluster. An Ingress can be configured to pick a specific one via `spec.ingressClassName`.

Ingress Traffic Routing

The context path is mapped to a backend (Service name and port)



Creating an Ingress with Imperative Approach

The rule definition requires intricate knowledge of syntax

```
$ kubectl create ingress corellian  
--rule="star-alliance.com/corellian/api=corellian:8080"  
ingress.networking.k8s.io/corellian created
```

Ingress YAML Manifest

Allows for defining multiple rules that map to backend

```
apiVersion: networking.k8s.io/v1
kind: Ingress
metadata:
  name: corellian
spec:
  rules:
  - host: star-alliance.com
    http:
      paths:
      - backend:
          service:
            name: corellian
            port:
              number: 8080
        path: /corellian/api
        pathType: Exact
```

← The host definition is optional

Ingress Rules

Traffic routing is controlled by rules defined on Ingress resource

Type	Example	Description
An optional host	<code>mycompany.abc.com</code>	If a host is provided, then the rules apply to that host. If no host is defined, then all inbound HTTP(S) traffic is handled.
A list of paths The	<code>/corellian/api</code>	Incoming traffic must match the host and path to correctly forward the traffic to a Service.
backend	<code>corellian:8080</code>	A combination of Service name and port.

Path Types

Incoming requests match based on assigned type in rule

Path Type	Rule	Incoming Request
Exact	/corellian/api	Matches /corellian/api but does not match /corellian/test or /corellian/api/.
Prefix	/corellian/api	Matches /corellian/api and /corellian/api/ but does not match /corellian/test.

Listing and Describing Ingresses

Renders hosts, IP addresses, and ports

```
$ kubectl get ingress
```

NAME	CLASS	HOSTS	ADDRESS	PORTS	AGE
corellian	<none>	star-alliance.com	192.168.64.15	80	10m

```
$ kubectl describe ingress corellian
```

```
Name:                corellian
Namespace:           default
Address:             192.168.64.15
Default backend:     default-http-backend:80 (<error: \
                    endpoints "default-http-backend" not found>)

Rules:
  Host                Path  Backends
  ----                -
  star-alliance.com  /corellian/api  corellian:8080 (172.17.0.5:8080)

Annotations:         <none>
Events:              <none>
```

Configuring DNS for an Ingress

Making an Ingress convenient to use by end users

- To resolve the Ingress, you'll need to configure DNS entries to the external address.
- You'll need to either configure an A record, or a CNAME record. The [ExternalDNS](#) cluster add-on can help with managing those DNS records for you.
- If you have no domain or are using a local Kubernetes solution, you can add the mapping between IP address and domain name to `/etc/hosts`.

Accessing an Ingress

Use host name, port, and context path

```
$ wget star-alliance.com/coreellian/api --timeout=5 --tries=1
--2021-11-30 19:34:57-- http://star-alliance.com/coreellian/api
Resolving star-alliance.com (star-alliance.com)... 192.168.64.15
Connecting to star-alliance.com (star-alliance.com)|192.168.64.15|:80... ↵
connected.
HTTP request sent, awaiting response... 200 OK
...

$ wget star-alliance.com/coreellian/api/ --timeout=5 --tries=1
--2021-11-30 15:36:26-- http://star-alliance.com/coreellian/api/
Resolving star-alliance.com (star-alliance.com)... 192.168.64.15
Connecting to star-alliance.com (star-alliance.com)|192.168.64.15|:80... ↵
connected.
HTTP request sent, awaiting response... 404 Not Found
2021-11-30 15:36:26 ERROR 404: Not Found.
```

Lab 19

Defining and using an Ingress





Network Policies

Restricting Pod-to-Pod communication

What is a Network Policy?

Firewall rules for Pod-to-Pod communication

- The installed [Container Network Interface \(CNI\) plugin](#) leases and assigns a new virtual IP address to every Pod when it is created.
- Communication between Pods within the same cluster is unrestricted. You can use a Pod's IP address to communicate with it.
- Network policies implement the principle of least privilege. Only allow Pods to communicate if it is needed to fulfill the requirements of architecture.
- You can disallow and allow network communication with coarse-grained and fine-grained rules.

What is an Network Policy Controller?

Evaluating Network Policy rules

- A Network Policy cannot work without a Network Policy controller. The Network Policy controller evaluates the collection of rules defined by a Network Policy.
- One option for such a Network Policy controller is [Cilium](#).

Network Policy YAML Manifest

Key-value pairs can be parsed from different sources

```
apiVersion: networking.k8s.io/v1
kind: NetworkPolicy
metadata:
  name: api-allow
spec:
  podSelector:
    matchLabels:
      app: payment-processor
      role: api
  ingress:
    - from:
        - podSelector:
            matchLabels:
              app: coffeeshop
```

Selects the Pod the policy should apply to by label selection

Allows incoming traffic from the Pod with matching labels within the same namespace

Network Policy Rules

Attributes that form the rules

Attribute	Description
<code>podSelector</code>	Selects the Pods in the namespace to apply the network policy to.
<code>policyTypes</code>	Defines the type of traffic (i.e., ingress and/or egress) the network policy applies to.
<code>ingress</code>	Lists the rules for incoming traffic. Each rule can define <code>from</code> and <code>ports</code> sections.
<code>egress</code>	Lists the rules for outgoing traffic. Each rule can define <code>to</code> and <code>ports</code> sections.


Behavior of `to` and `from` Selectors

Ingress and egress sections define the same attributes

Attribute	Description
<code>podSelector</code>	Selects Pods by label(s) in the same namespace as the Network Policy which should be allowed as ingress sources or egress destinations.
<code>namespaceSelector</code>	Selects namespaces by label(s) for which all Pods should be allowed as ingress sources or egress destinations.
<code>namespaceSelector</code> and <code>podSelector</code>	Selects Pods by label(s) within namespaces by label(s).

Example Network Policy

Restricting access to the payment processor Pod from other Pods



Listing Network Policies

Doesn't provide much details apart from Pod selector labels

```
$ kubectl get networkpolicy
```

NAME	POD-SELECTOR	AGE
api-allow	app=payment-processor,role=api	83m

Rendering Network Policy Details

Ingress and egress rules can be inspected in details

```
$ kubectl describe networkpolicy api-allow
Name:          api-allow
Namespace:     default
Created on:    2020-09-26 18:02:57 -0600 MDT
Labels:        <none>
Annotations:   <none>
Spec:
  PodSelector:    app=payment-processor,role=api
  Allowing ingress traffic:
    To Port: <any> (traffic allowed to all ports)
    From:
      PodSelector: app=coffeeshop
  Not affecting egress traffic
  Policy Types: Ingress
```

Verifying the Runtime Behavior

Communication from Pods not matching with labels will be blocked

```
$ kubectl run grocery-store --rm -it --image=busybox ↵  
-l app=grocery-store,role=backend -- /bin/sh  
/ # wget --spider --timeout=1 10.0.0.51  
Connecting to 10.0.0.51 (10.0.0.51:80)  
wget: download timed out  
/ # exit  
pod "grocery-store" deleted
```

```
$ kubectl run coffeeshop --rm -it --image=busybox ↵  
-l app=coffeeshop,role=backend -- /bin/sh  
/ # wget --spider --timeout=1 10.0.0.51  
Connecting to 10.0.0.51 (10.0.0.51:80)  
remote file exists  
/ # exit  
pod "coffeeshop" deleted
```

Default Deny Network Policies

Network policies are additive

```
apiVersion: networking.k8s.io/v1
kind: NetworkPolicy
metadata:
  name: default-deny-all
spec:
  podSelector: {}
  policyTypes:
    - Ingress
    - Egress
```

Selects all Pods in
the namespace

Applies in incoming
and outgoing traffic

Restricting Access to Ports

By default, all ports are accessible

```
apiVersion: networking.k8s.io/v1
kind: NetworkPolicy
metadata:
  name: port-allow
spec:
  podSelector:
    matchLabels:
      app: backend
  ingress:
    - from:
        - podSelector:
            matchLabels:
              app: frontend
      ports:
        - protocol: TCP
          port: 8080
```

Only allow incoming
traffic to port 8080

Lab 20

Restricting access to and from a Pod with Network Policies

